



TEWA 1: Advanced Data Analysis

Lecture II

Lei Zhang

Social, Cognitive and Affective Neuroscience Unit (SCAN-Unit)
Department of Cognition, Emotion, and Methods in Psychology

https://github.com/lei-zhang/tewa1_univie

lei.zhang@univie.ac.at
lei-zhang.net
@lei_zhang_lz



universität
wien

Fakultät für Psychologie


Bayesian warm-up?

Hierarchical Structure




cognitive model

statistics



computing

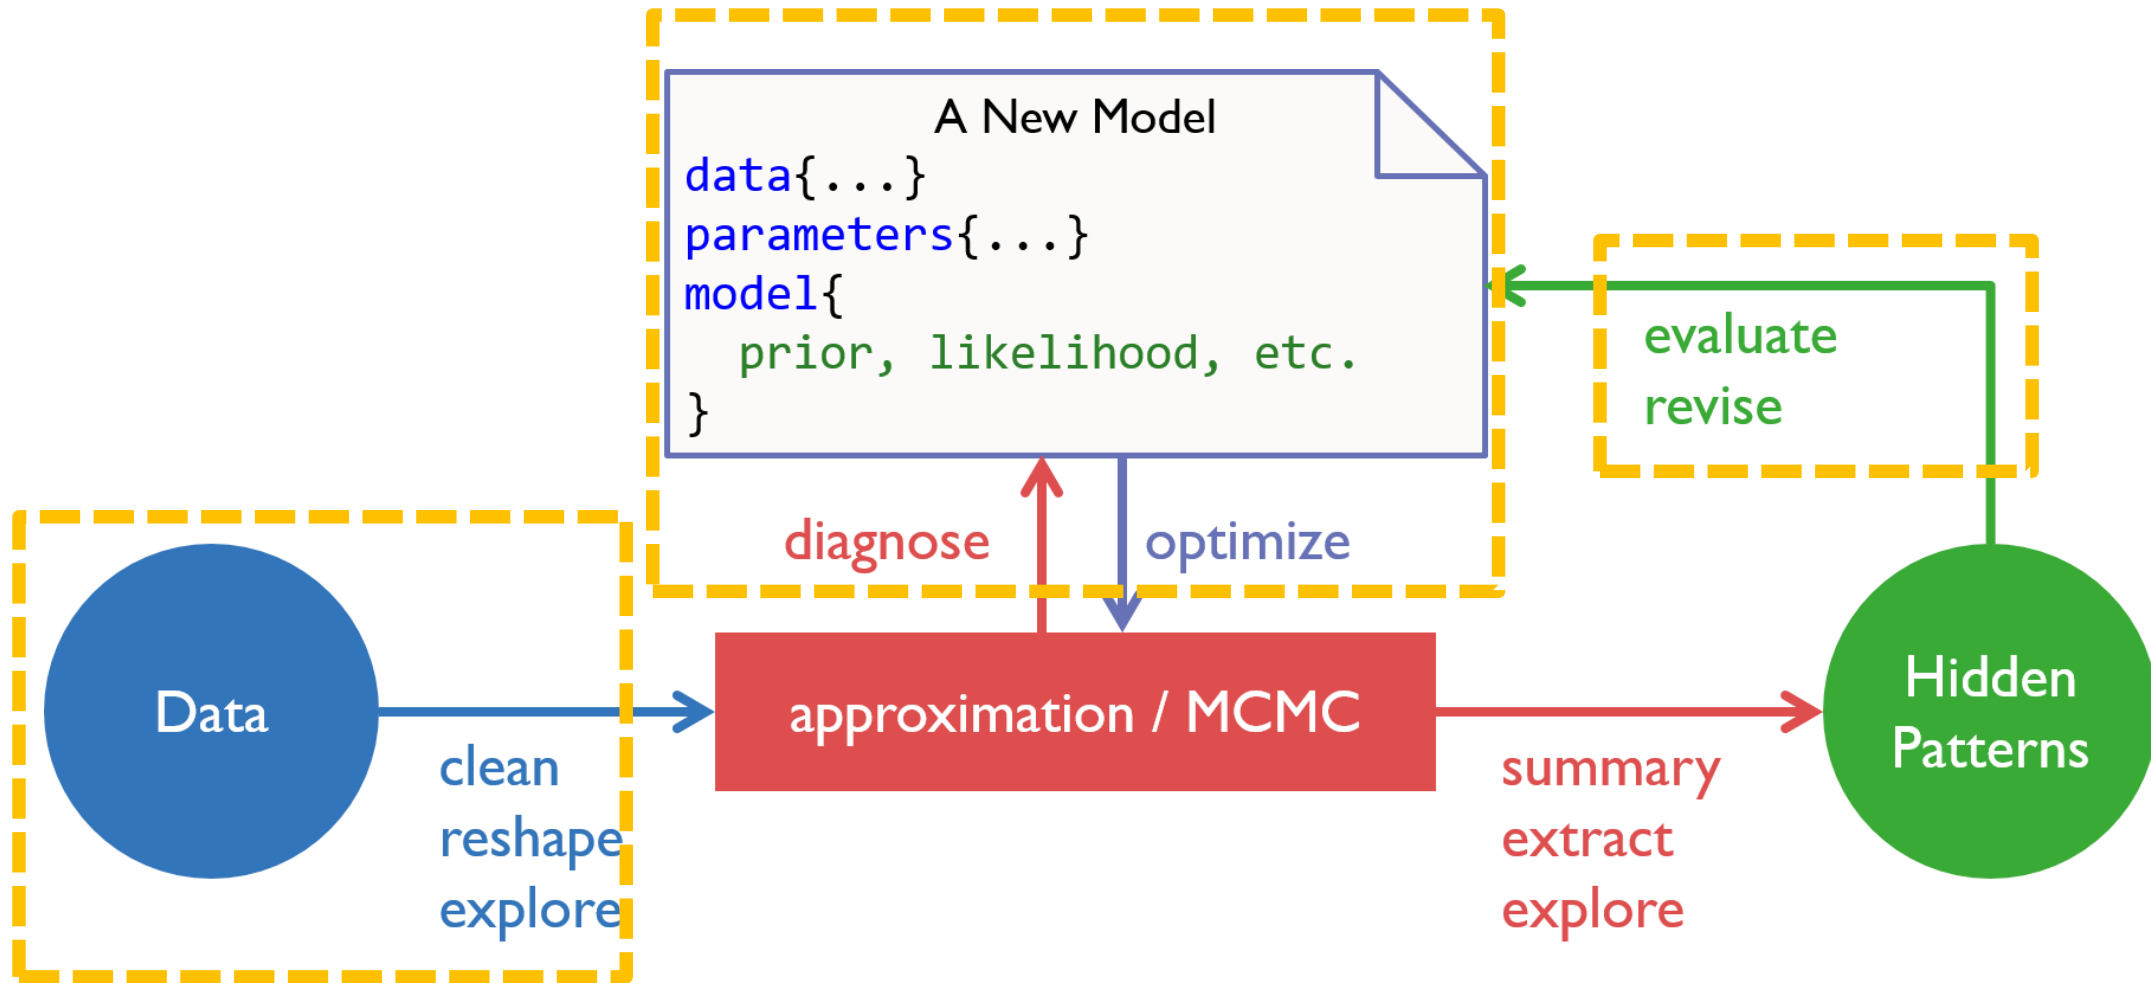


$$P(\Theta, \Phi | D) = \frac{P(D | \Theta, \Phi) P(\Theta, \Phi)}{P(D)} \propto P(D | \Theta) P(\Theta | \Phi) P(\Phi)$$



OPTIMIZING STAN CODES





Optimizing Stan Code



cognitive model

statistics

computing

Preprocess data

run as many calculations as you can outside Stan

Specify a proper model

follow literature, supervision, experience, etc.

Vectorizing

vectorize Stan code whenever you can

Reparameterizing

reparameterize target parameter to simple distributions

Preprocess Data

cognitive model

statistics

computing

$$\overline{\text{height}} = \alpha + \beta_1 * \text{weight} + \beta_2 * \text{weight}^2$$

```
d$weight_sq <- d$weight^2
```

```
data {  
  int<lower=0> N;  
  vector<lower=0>[N] height;  
  vector<lower=0>[N] weight;  
  vector<lower=0>[N] weight_sq;  
}
```

Specify a Proper Model

cognitive model

statistics

computing

- Visualize your data
- Follow Literatures
- Start from simple and then build complexities
- Simulate data and run model recovery

A New Model

```
data{...}  
parameters{...}  
model{  
  prior, likelihood, etc.  
}
```


Vectorization

cognitive model

statistics

computing

```
model {  
  for (n in 1:N) {  
    flip[n] ~ bernoulli(theta);  
  }  
}
```



```
model {  
  flip ~ bernoulli(theta);  
}
```

```
parameters {  
  ...  
  real<lower=0,upper=1> lr[nSubjects];  
  real<lower=0,upper=3> tau[nSubjects];  
}  
  
model {  
  ...  
  lr ~ normal(lr_mu, lr_sd) ;  
  tau ~ normal(tau_mu, tau_sd) ;  
  ...  
}
```

```
model {  
  vector[N] mu;  
  for (i in 1:N) {  
    mu[i] = alpha + beta * weight[i];  
    height[i] ~ normal(mu[i], sigma)  
  }  
}
```



```
model {  
  vector[N] mu;  
  mu = alpha + beta * weight;  
  height ~ normal(mu, sigma);  
}
```



```
model {  
  height ~ normal(alpha + beta * weight, sigma);  
}
```

Reparameterization

Neal's Funnel


cognitive model

statistics

computing

$$p(y, x) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```




Non-centered Reparameterization*

cognitive model

statistics

computing


$$\theta \sim \text{Normal}(\mu_\theta, \sigma_\theta)$$



$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \mu_\theta + \sigma_\theta \tilde{\theta}$$

Stan likes **simple**
distributions!



Reparameterization


Neal's Funnel

$$p(y, x) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```



```
parameters {  
  real y_raw;  
  vector[9] x_raw;  
}  
transformed parameters {  
  real y;  
  vector[9] x;  
  
  y = 3.0 * y_raw;  
  x = exp(y/2) * x_raw;  
}  
model {  
  y_raw ~ normal(0, 1);  
  x_raw ~ normal(0, 1);  
}
```



Stan Sampling Parameters

cognitive model

statistics

computing

parameter	description	constraint	default
iterations	number of MCMC samples (per chain)	int, > 0	2000
delta: δ	target Metropolis acceptance rate	$\delta \in [0, 1]$	0.80
stepsize: ε	initial HMC step size	real, $\varepsilon > 0$	2.0
max_treedepth: L	maximum HMC steps per iteration	int, $L > 0$	10

Typical adjustments

- Increase iterations
- Increase delta
- Decrease stepsize
- Might have to increase max_treedepth

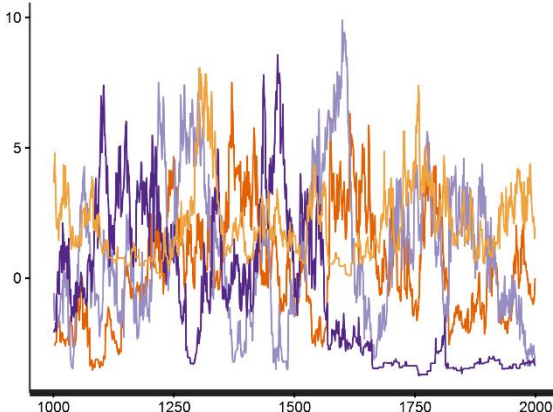
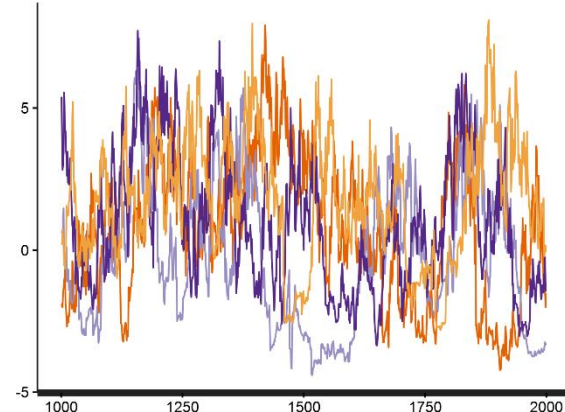
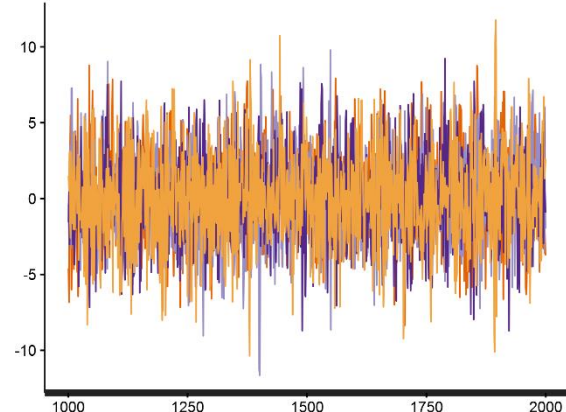
```
funnel_fit2 <- stan("_scripts/funnel.stan",  
  iter = 4000,  
  control = list(adapt_delta = 0.999,  
    stepsize = 1.0,  
    max_treedepth = 20))
```

Neal's Funnel: Comparing Performance

cognitive model

statistics

computing

	direct model	adjusted direct model	reparameterized model
Rhat (y)	1.22	1.1	1.0
n_eff (y)	18	42	3886
runtime*	48.50 sec	50.76 sec	50.12 sec
n_eff (y) / runtime	0.37 / sec	0.82 / sec	77.53 / sec
n_divergent	53	0	0
traceplot (y)			

*: 2 cores in parallel, including compiling time

How about **Bounded** Parameters?

$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \mu_{\theta} + \sigma_{\theta} \tilde{\theta}$$

$$\theta \in (-\infty, +\infty)$$



$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta})$$

$$\theta \in [0, 1]$$

constraint	reparameterization
$\theta \in (-\infty, +\infty)$	$\theta = \mu_{\theta} + \sigma_{\theta} \tilde{\theta}$
$\theta \in [0, N]$	$\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}) \times N$
$\theta \in [M, N]$	$\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}) \times (N-M) + M$
$\theta \in (0, +\infty)$	$\theta = \exp(\mu_{\theta} + \sigma_{\theta} \tilde{\theta})$

Apply to Our Hierarchical RL Model

cognitive model

statistics

computing

```
parameters {  
  real<lower=0,upper=1> lr_mu;  
  real<lower=0,upper=3> tau_mu;  
  
  real<lower=0> lr_sd;  
  real<lower=0> tau_sd;  
  
  real<lower=0,upper=1> lr[nSubjects];  
  real<lower=0,upper=3> tau[nSubjects];  
}
```



```
parameters {  
  # group-level parameters  
  real lr_mu_raw;  
  real tau_mu_raw;  
  real<lower=0> lr_sd_raw;  
  real<lower=0> tau_sd_raw;  
  
  # subject-level raw parameters  
  vector[nSubjects] lr_raw;  
  vector[nSubjects] tau_raw;  
}  
  
transformed parameters {  
  vector<lower=0,upper=1>[nSubjects] lr;  
  vector<lower=0,upper=3>[nSubjects] tau;  
  
  for (s in 1:nSubjects) {  
    lr[s] = Phi_approx( lr_mu_raw + lr_sd_raw * lr_raw[s] );  
    tau[s] = Phi_approx( tau_mu_raw + tau_sd_raw * tau_raw[s] ) * 3;  
  }  
}
```


Apply to Our Hierarchical RL Model

cognitive model

statistics

computing

```
model {  
  lr_sd ~ cauchy(0,1);  
  tau_sd ~ cauchy(0,3);  
  lr ~ normal(lr_mu, lr_sd) ;  
  tau ~ normal(tau_mu, tau_sd) ;  
  
  for (s in 1:nSubjects) {  
    vector[2] v;  
    real pe;  
    v = initV;  
  
    for (t in 1:nTrials) {  
      choice[s,t] ~ categorical_logit( tau[s] * v );  
      pe = reward[s,t] - v[choice[s,t]];  
      v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;  
    }  
  }  
}
```



```
model {  
  lr_mu_raw ~ normal(0,1);  
  tau_mu_raw ~ normal(0,1);  
  lr_sd_raw ~ cauchy(0,3);  
  tau_sd_raw ~ cauchy(0,3);  
  
  lr_raw ~ normal(0,1);  
  tau_raw ~ normal(0,1);  
  
  for (s in 1:nSubjects) {  
    ...  
  }  
  
  generated quantities {  
    real<lower=0,upper=1> lr_mu;  
    real<lower=0,upper=3> tau_mu;  
  
    lr_mu = Phi_approx(lr_mu_raw);  
    tau_mu = Phi_approx(tau_mu_raw) * 3;  
  }  
}
```

Exercise XII

cognitive model

statistics

computing

```
.../07.optm_rl/_scripts/reinforcement_learning_hrch_main.R
```

TASK: (1) Complete the Matt Trick
(2) fit the optimized hierarchical RL model

```
> source('_scripts/reinforcement_learning_hrch_main.R')  
> fit_rl4 <- run_rl_mp2(optimized = TRUE)
```



Hierarchical Fitting – Optimized

cognitive model



statistics

computing

Posterior Means (hrch)



Posterior Means (hrch + optm)




Comparing with True Parameters

cognitive model


statistics

computing


Posterior Means (indy)






Posterior Means (hrch)



Posterior Means (hrch+optm)



True Parameters





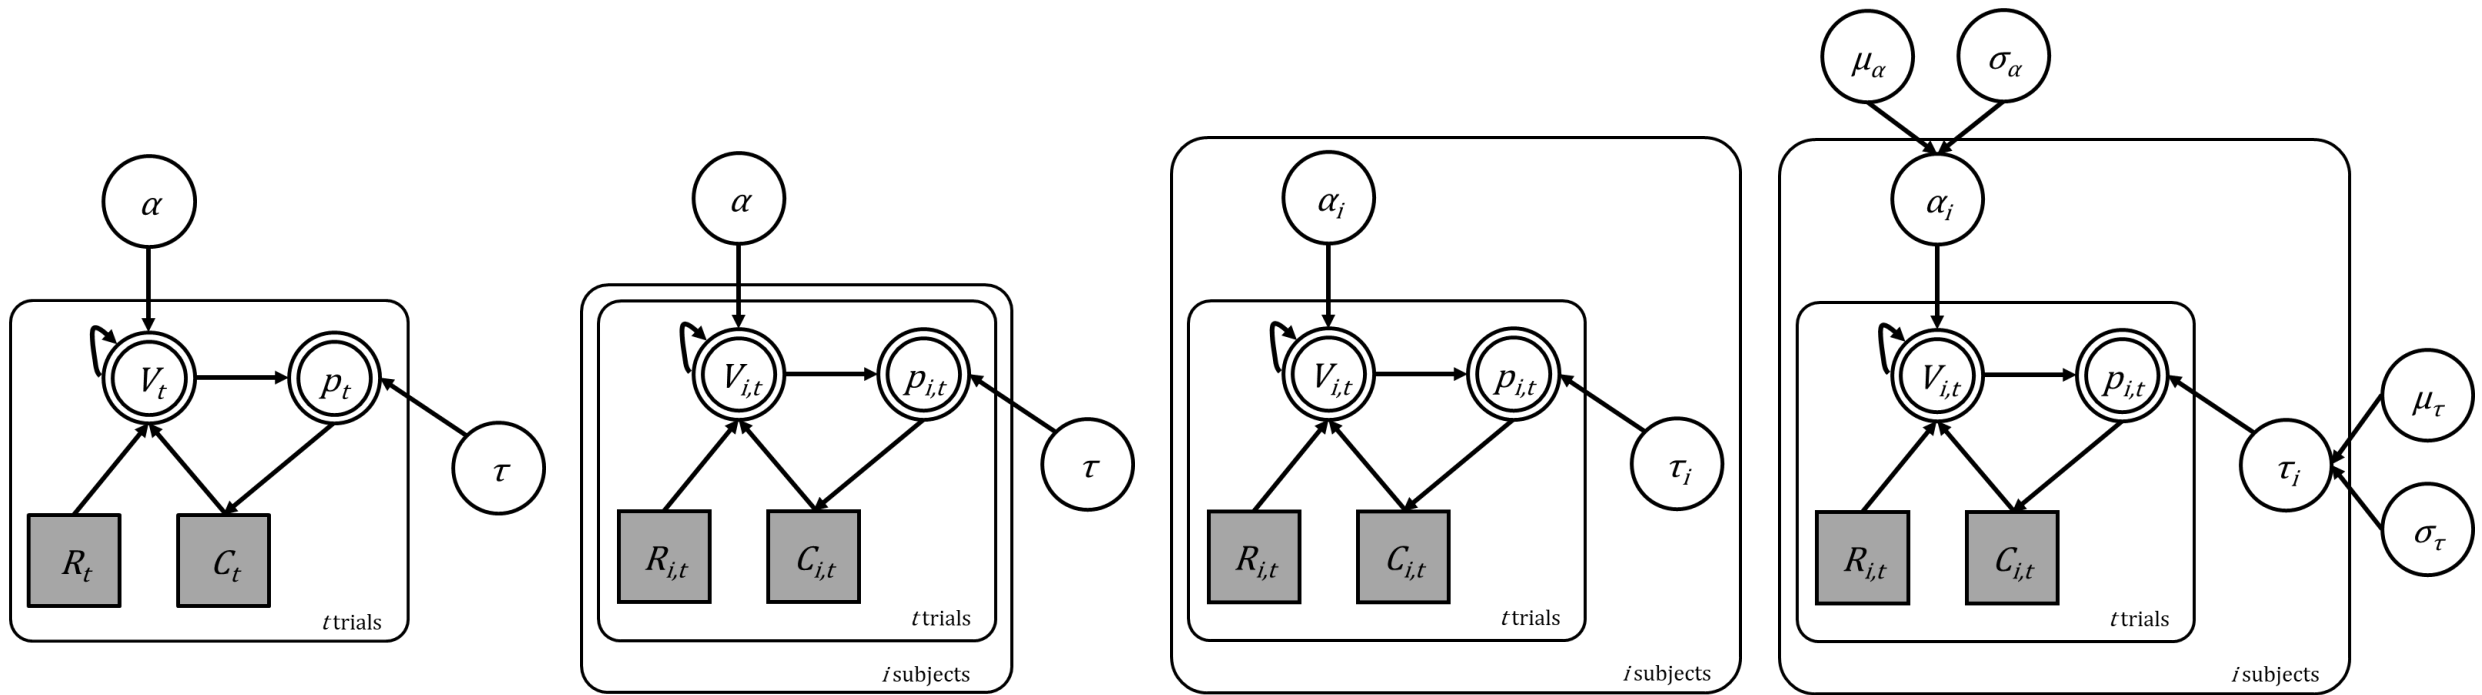
Posterior Predictive Check

cognitive model

statistics



computing

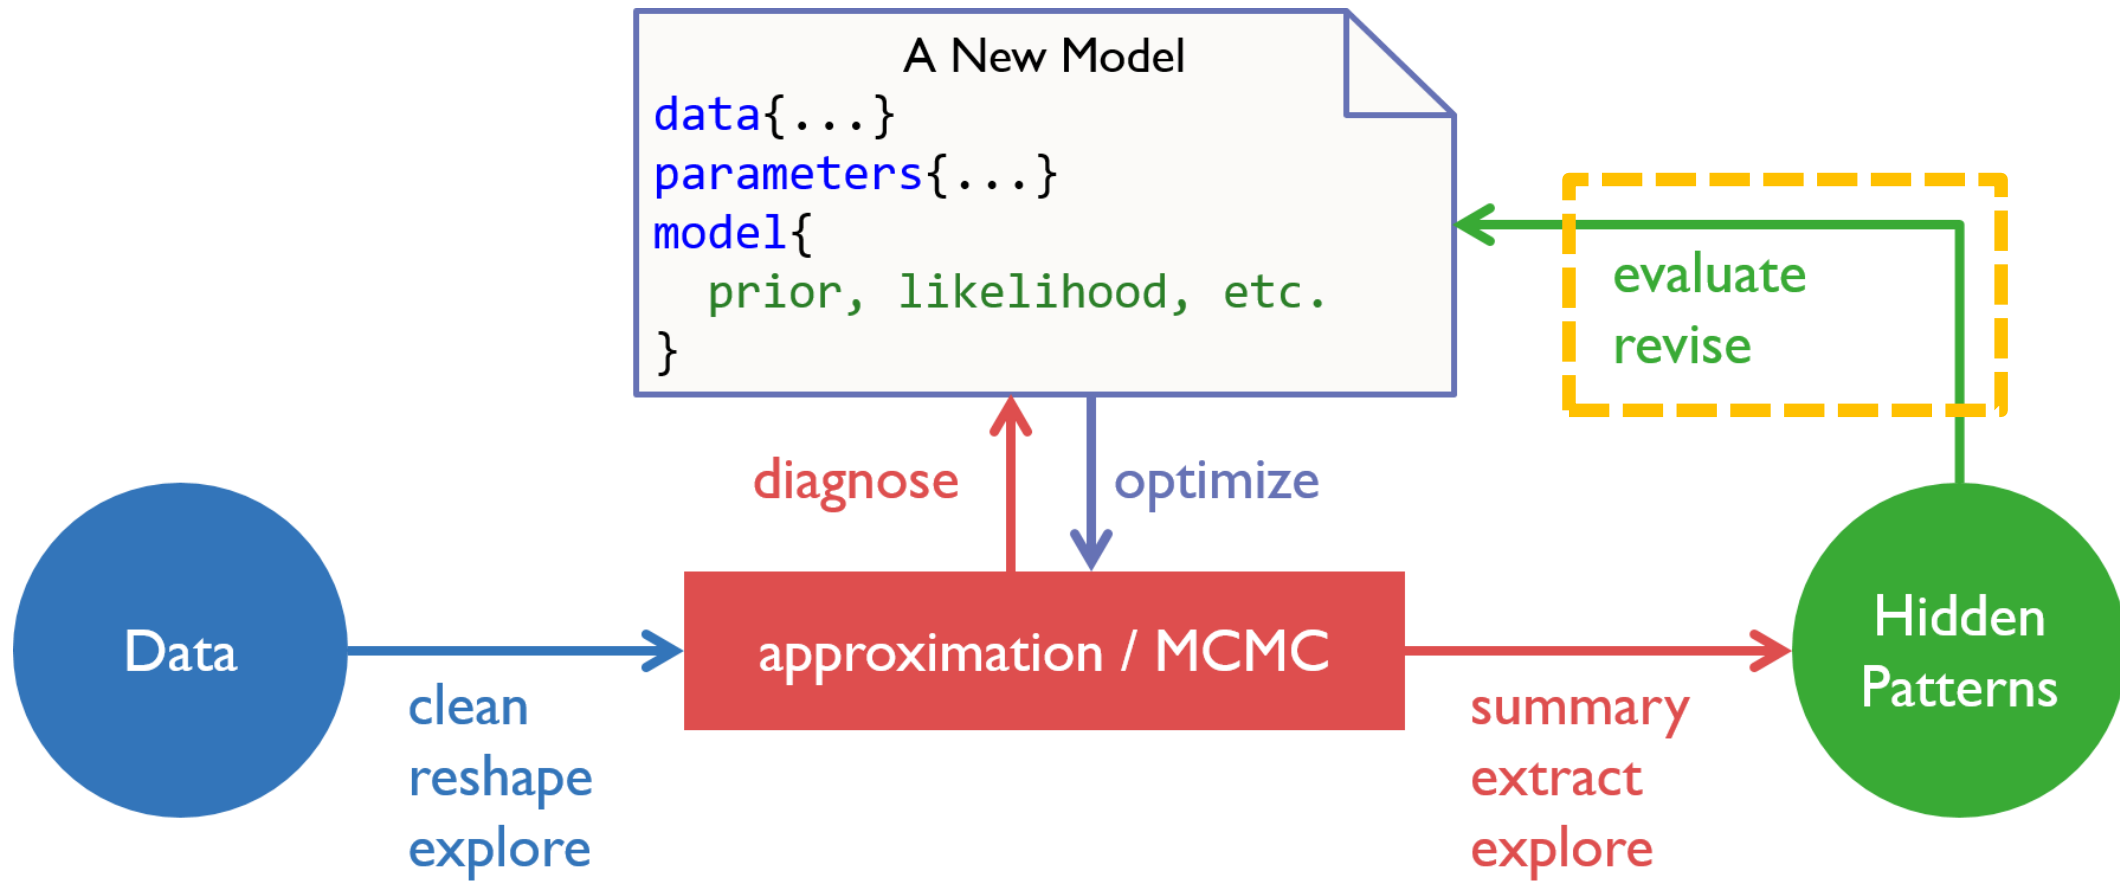




constraint	reparameterization
$\theta \in (-\infty, +\infty)$	$\theta = \mu_\theta + \sigma_\theta \tilde{\theta}$
$\theta \in [0, N]$	$\theta = \text{Probit}^{-1}(\mu_\theta + \sigma_\theta \tilde{\theta}) \times N$
$\theta \in [M, N]$	$\theta = \text{Probit}^{-1}(\mu_\theta + \sigma_\theta \tilde{\theta}) \times (N-M) + M$
$\theta \in (0, +\infty)$	$\theta = \exp(\mu_\theta + \sigma_\theta \tilde{\theta})$

MODEL COMPARISON





Model Comparison

cognitive model

statistics

computing


Which model provides the best **fit**?



Which model represents the best **balance** between model fit and model complexity?

Ockham's razor:

Models with fewer assumptions are to be preferred



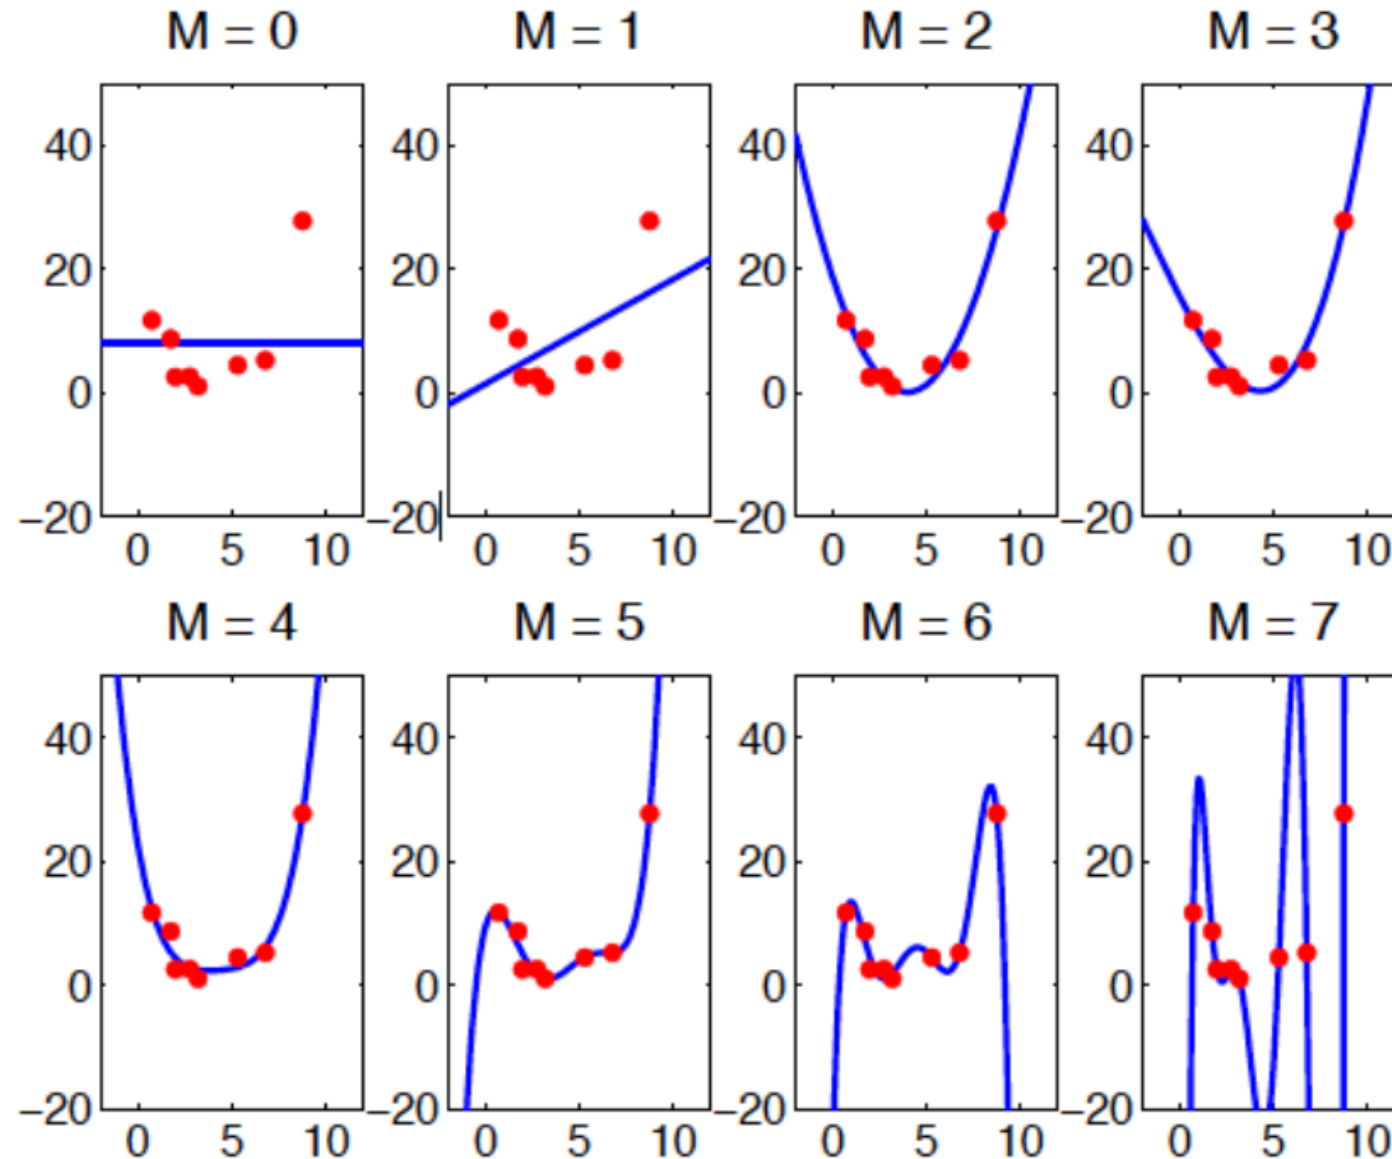
- overfitting: learn **too much** from the data
- underfitting: learn **too little** from the data

Which model has the highest predictive power?

cognitive model

statistics

computing




Focusing on Predictive Accuracy

cognitive model

statistics

computing



- Nothing prevents you from doing that in a Bayesian context but holding out data makes your posterior distribution more diffuse
- Bayesians usually condition on *all* the data and evaluate how well a model is expected to **predict out of sample** using "information criteria": model with the **highest expected log predictive density (ELPD)** for new data

Information Criteria

cognitive model

statistics

computing

AIC – Akaike information criterion

DIC – Deviance Information Criterion

WAIC – Widely Applicable Information Criterion

finding the model that has
the highest out-of-sample
predictive accuracy

approximation to LOO

BIC – Bayesian Information Criterion

finding the “true” model

Compute WAIC from Likelihood

cognitive model

statistics

computing

$$\text{WAIC} = -2 \widehat{\text{elpd}}_{\text{waic}}$$

expected log pointwise predictive density

$$\widehat{\text{elpd}}_{\text{waic}} = \widehat{\text{lpd}} - \widehat{p}_{\text{waic}}$$

$\widehat{\text{lpd}}$ = computed log pointwise predictive density

$$= \sum_{i=1}^n \log \left(\frac{1}{S} \sum_{s=1}^S p(y_i | \theta^s) \right).$$

```
lpd <- log(colMeans(exp(log_lik)))
```

estimated effective number of parameters

$$\widehat{p}_{\text{waic}} = \sum_{i=1}^n V_{s=1}^S (\log p(y_i | \theta^s))$$

```
p_waic <- colVars(log_lik)
```

*IC comparisons

cognitive model

statistics

computing

		No pooling ($\tau = \infty$)	Complete pooling ($\tau = 0$)	Hierarchical model (τ estimated)
AIC	$-2 \text{ lpd} = -2 \log p(y \hat{\theta}_{\text{mle}})$	54.6	59.4	
	k	8.0	1.0	
	$\text{AIC} = -2 \widehat{\text{elpd}}_{\text{AIC}}$	70.6	61.4	
DIC	$-2 \text{ lpd} = -2 \log p(y \hat{\theta}_{\text{Bayes}})$	54.6	59.4	57.4
	p_{DIC}	8.0	1.0	2.8
	$\text{DIC} = -2 \widehat{\text{elpd}}_{\text{DIC}}$	70.6	61.4	63.0
WAIC	$-2 \text{ lppd} = -2 \sum_i \log p_{\text{post}}(y_i)$	60.2	59.8	59.2
	$p_{\text{WAIC } 1}$	2.5	0.6	1.0
	$p_{\text{WAIC } 2}$	4.0	0.7	1.3
	$\text{WAIC} = -2 \widehat{\text{elppd}}_{\text{WAIC } 2}$	68.2	61.2	61.8
LOO-CV	-2 lppd		59.8	59.2
	$p_{\text{loo-cv}}$		0.5	1.8
	$-2 \text{ lppd}_{\text{loo-cv}}$		60.8	62.8

Recording the Log-Likelihood in Stan

cognitive model

statistics

computing

```
generated quantities {  
  ...  
  real log_lik[nSubjects];  
  ...  
  
  { # local section, this saves time and space  
    for (s in 1:nSubjects) {  
      vector[2] v;  
      real pe;  
  
      log_lik[s] = 0;  
      v = initV;  
  
      for (t in 1:nTrials) {  
        log_lik[s] = log_lik[s] + categorical_logit_lpmf(choice[s,t] | tau[s] * v);  
  
        pe = reward[s,t] - v[choice[s,t]];  
        v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;  
      }  
    }  
  }  
}
```

The {loo} Package

cognitive model

statistics

computing

```
> library(loo)
> LL1    <- extract_log_lik(stanfit)
> loo1   <- loo(LL1)    # PSIS leave-one-out
> waic1  <- waic(LL1)   # WAIC
```

Computed from 4000 by 20 log-likelihood matrix

	Estimate	SE
elpd_loo	-29.5	3.3
p_loo	2.7	1.0
looic	58.9	6.7




Pareto Smoothed Importance Sampling

Reversal Learning Task

cognitive model

statistics

computing



Fictitious RL
(Counterfactual RL)

Value update:

$$V_{t+1}^c = V_t^c + \alpha * PE$$

$$V_{t+1}^{nc} = V_t^{nc} + \alpha * PEnc$$

Prediction error:

$$PE = R_t - V_t^c$$




$$PEnc = -R_t - V_t^{nc}$$

More on Fictitious RL

cognitive model

statistics

computing



Exercise XIII

cognitive model

statistics

computing

```
.../08.compare_models/_scripts/compare_models_main.R
```

TASK: (1) complete the fictitious RL model (model2, loglik)
(2) fit and compare the 2 models

Exercise XIII – output

```
> LL1 <- extract_log_lik(fit_rl1)
```

```
> ( loo1 <- loo(LL1) )
```

Computed from 4000 by 10 log-likelihood matrix

	Estimate	SE
elpd_loo	-389.8	15.4
p_loo	3.8	0.8
looic	779.5	30.9

```
> ( loo2 <- loo(LL2) )
```

Computed from 4000 by 10 log-likelihood matrix



	Estimate	SE
elpd_loo	-281.3	17.5
p_loo	3.4	0.5
looic	562.6	35.0

Posterior Predictive Check

cognitive model

statistics

computing



ANY
QUESTIONS
?

Happy Computing!